

Problem 2

- In the case described we are using a linear function $W \in R^{K \times D}$ to map data in from R^D into R^K . In particular requiring a zero reconstruction error would require

$$BW X = X \iff BW = Id$$

Where $B \in R^{D \times K}$. But as the $rank(W)$ is K and not D , we do not have existence of a left inverse.

- It is possible if $K \geq D$ because we then can define a left inverse.

Problem 3

Compute the expected value:

$$E[X] = E\left[\sum_k \pi_k N(x|\mu_k, \Sigma_k)\right] = \sum_k E[\pi_k N(x|\mu_k, \Sigma_k)] = \sum_k \pi_k E[N(x|\mu_k, \Sigma_k)] = \sum_k \pi_k \mu_k$$

Compute the covariance:

$$Cov[X] = E[xx^T] - E[x]E[x]^T = \sum_k \pi_k E[xx^T] - E[x]E[x]^T$$

We know that $E[xx^T] = Cov[x] + E[x]E[x]^T = \sum_k \pi_k (\Sigma_k + \mu_k \mu_k^T)$

We can conclude that:

$$Cov[X] = \sum_k \pi_k (\Sigma_k + \mu_k \mu_k^T) - E[X]E[X]^T$$

Problem 4

a)

- First we sample the cluster k from $Cat(\pi^x)$
- We draw the sample x from the cluster k : $x \sim \mathcal{N}(\mu_k^x, \Sigma_k^x)$
- We sample the cluster l from $Cat(\pi^y)$
- We draw the sample y from cluster l : $y \sim \mathcal{N}(\mu_l^y, \Sigma_l^y)$
- We can compute $z = y + x = \mathcal{N}(\mu_l^y, \Sigma_l^y) + \mathcal{N}(\mu_k^x, \Sigma_k^x)$

b) In a) we computed $z = y + x = \mathcal{N}(\mu_l^x, \Sigma_l^x) + \mathcal{N}(\mu_k^x, \Sigma_k^x)$, z is a gaussian because it is the sum of two Gaussian distributions: $z \sim \mathcal{N}(\mu_l^y + \mu_k^x, \Sigma_l^y + \Sigma_k^x)$. The probability of the sample from $\mathcal{N}(\mu_l^y + \mu_k^x, \Sigma_l^y + \Sigma_k^x)$ is $\pi_l^y * \pi_k^x$. To compute the likelihood of z we have a summation of the product of a probability π and a Gaussian so the result is again a Gaussian mixture model.

c)

$$P(z|\pi_z, \mu_z, \Sigma_z) = \sum_{k=1} P(C_k = 1|\pi) P(z|C_k = 1, \mu_k^z, \Sigma_k^z) = \sum_{k=1}^K \sum_{l=1}^K \pi_l^x \pi_k^y \mathcal{N}(z|\mu_k^x + \mu_l^y, \Sigma_k^x + \Sigma_l^y)$$

Exporting the results to PDF

Once you complete the assignments, export the entire notebook as PDF and attach it to your homework solutions. The best way of doing that is

1. Run all the cells of the notebook.
2. Export/download the notebook as PDF (File -> Download as -> PDF via LaTeX (.pdf)).
3. Concatenate your solutions for other tasks with the output of Step 2. On linux, you can use `pdffunite`, there are similar tools for other platforms, too. You can only upload a single PDF file to Moodle.

Make sure you are using `nbconvert` version 5.5 or later by running `jupyter nbconvert --version`. Older versions clip lines that exceed page width, which makes your code harder to grade.

Matrix Factorization

In [1]:

```
import time
import scipy.sparse as sp
import numpy as np
from scipy.sparse.linalg import svds
from sklearn.linear_model import Ridge

import matplotlib.pyplot as plt
%matplotlib inline
```

Restaurant recommendation

The goal of this task is to recommend restaurants to users based on the rating data in the Yelp dataset. For this, we try to predict the rating a user will give to a restaurant they have not yet rated based on a latent factor model.

Specifically, the objective function (loss) we wanted to optimize is:

$$\mathcal{L} = \min_{P,Q} \sum_{(i,x) \in W} (M_{ix} - \mathbf{q}_i^T \mathbf{p}_x)^2 + \lambda \sum_x \|\mathbf{p}_x\|^2 + \lambda \sum_i \|\mathbf{q}_i\|^2$$

where W is the set of (i, x) pairs for which the rating M_{ix} given by user i to restaurant x is known. Here we have also introduced two regularization terms to help us with overfitting where λ is hyper-parameter that control the strength of the regularization.

Hint 1: Using the closed form solution for regression might lead to singular values. To avoid this issue perform the regression step with an existing package such as scikit-learn. It is advisable to use ridge regression to account for regularization.

Hint 2: If you are using the scikit-learn package remember to set `fit_intercept = False` to only learn the coefficients of the linear regression.

Load and Preprocess the Data (nothing to do here)

In [2]:

```
ratings = np.load("exercise_12_matrix_factorization_ratings.npy")
```

In [3]:

```
# We have triplets of (user, restaurant, rating).
ratings
```

Out[3]:

```
array([[101968,    1880,     1],
       [101968,     284,     5],
       [101968,   1378,     2],
       ...,
       [ 72452,    2100,     4],
       [ 72452,    2050,     5],
       [ 74861,    3979,     5]])
```

Now we transform the data into a matrix of dimension $[N, D]$, where N is the number of users and D is the number of restaurants in the dataset. We store the data as a sparse matrix to avoid out-of-memory issues.

In [4]:

```
n_users = np.max(ratings[:,0] + 1)
n_restaurants = np.max(ratings[:,1] + 1)
M = sp.coo_matrix((ratings[:,2], (ratings[:,0], ratings[:,1])),
, shape=(n_users, n_restaurants)).tocsr()
M
```

Out[4]:

```
<337867x5899 sparse matrix of type '<class 'numpy.
longlong'>'
      with 929606 stored elements in Compressed
Sparse Row format>
```

To avoid the [cold start problem](https://en.wikipedia.org/wiki/Cold_start_(computing)) ([https://en.wikipedia.org/wiki/Cold_start_\(computing\)](https://en.wikipedia.org/wiki/Cold_start_(computing))), in the preprocessing step, we recursively remove all users and restaurants with 10 or less ratings.

Then, we randomly select 200 data points for the validation and test sets, respectively.

After this, we subtract the mean rating for each users to account for this global effect.

Note: Some entries might become zero in this process -- but these entries are different than the 'unknown' zeros in the matrix. We store the indices for which we the rating data available in a separate variable.

In [5]:

```
def cold_start_preprocessing(matrix, min_entries):
    """
    Recursively removes rows and columns from the input matrix
    which have less than min_entries nonzero entries.

    Parameters
    -----
    matrix      : sp.spmatrix, shape [N, D]
                  The input matrix to be preprocessed.
    min_entries : int
                  Minimum number of nonzero elements per row a
nd column.

    Returns
    -----
    matrix      : sp.spmatrix, shape [N', D']
                  The pre-processed matrix, where  $N' \leq N$  and
 $D' \leq D$ 

    """
    print("Shape before: {}".format(matrix.shape))

    shape = (-1, -1)
    while matrix.shape != shape:
        shape = matrix.shape
        nnz = matrix>0
        row_ixs = nnz.sum(1).A1 > min_entries
        matrix = matrix[row_ixs]
        nnz = matrix>0
        col_ixs = nnz.sum(0).A1 > min_entries
        matrix = matrix[:,col_ixs]
    print("Shape after: {}".format(matrix.shape))
    nnz = matrix>0
    assert (nnz.sum(0).A1 > min_entries).all()
    assert (nnz.sum(1).A1 > min_entries).all()
    return matrix
```

Task 1: Implement a function that subtracts the mean user rating from the sparse rating matrix

In [6]:

```
def shift_user_mean(matrix):
    """
    Subtract the mean rating per user from the non-zero elements in the input matrix.

    Parameters
    -----
    matrix : sp.spmatrix, shape [N, D]
        Input sparse matrix.
    Returns
    -----
    matrix : sp.spmatrix, shape [N, D]
        The modified input matrix.

    user_means : np.array, shape [N, 1]
        The mean rating per user that can be used to recover the absolute ratings from the mean-shifted ones.

    """

    # TODO: Compute the modified matrix and user_means
    sum_ = np.squeeze(np.asarray(matrix.sum(1)))
    nonZero = np.diff(matrix.tocsr().indptr)
    user_means = (sum_/nonZero)

    diag = sp.diags(user_means,0)

    matrix_copy = matrix.copy()
    matrix_copy.data = np.ones_like(matrix_copy.data)

    matrix = matrix - (diag*matrix_copy).todense()
    user_means = user_means.reshape(len(user_means), 1)
    user_means = np.asmatrix(user_means)
    assert np.all(np.isclose(matrix.mean(1), 0))
    return matrix, user_means
```

Split the data into a train, validation and test set (nothing to do here)

In [7]:

```
def split_data(matrix, n_validation, n_test):
    """
```


Extract validation and test entries from the input matrix.

Parameters

matrix : *sp.spmatrix, shape [N, D]*
The input data matrix.

n_validation : *int*
The number of validation entries to extract.

n_test : *int*
The number of test entries to extract.

Returns

matrix_split : *sp.spmatrix, shape [N, D]*
A copy of the input matrix in which the validation and test entries have been set to zero.

val_idx : *tuple, shape [2, n_validation]*
The indices of the validation entries.

test_idx : *tuple, shape [2, n_test]*
The indices of the test entries.

val_values : *np.array, shape [n_validation,]*
The values of the input matrix at the validation indices.

test_values : *np.array, shape [n_test,]*
The values of the input matrix at the test indices.

"""

```
matrix_cp = matrix.copy()
non_zero_idx = np.argwhere(matrix_cp)
ixs = np.random.permutation(non_zero_idx)
val_idx = tuple(ixs[:n_validation].T)
test_idx = tuple(ixs[n_validation:n_validation + n_test].T
)
```

val_values = matrix_cp[val_idx].A1
test_values = matrix_cp[test_idx].A1

matrix_cp[val_idx] = matrix_cp[test_idx] = 0
matrix_cp.eliminate_zeros()

return matrix_cp, val_idx, test_idx, val_values, test_values

```
return matrix_cp, val_idx, test_idx, val_values, test_valu  
es
```

In [8]:

```
M = cold_start_preprocessing(M, 20)
```

Shape before: (337867, 5899)

Shape after: (3529, 2072)

In [9]:

```
n_validation = 200  
n_test = 200  
# Split data  
M_train, val_idx, test_idx, val_values, test_values = split_da  
ta(M, n_validation, n_test)
```

In [10]:

```
# Remove user means.  
nonzero_indices = np.argwhere(M_train)  
M_shifted, user_means = shift_user_mean(M_train)  
# Apply the same shift to the validation and test data.  
val_values_shifted = val_values - user_means[np.array(val_idx)  
.T[:,0]].A1  
test_values_shifted = test_values - user_means[np.array(test_i  
dx).T[:,0]].A1
```

Compute the loss function (nothing to do here)

In [11]:

```
def loss(values, ixs, Q, P, reg_lambda):
    """
    Compute the loss of the latent factor model (at indices ix
    s).

    Parameters
    -----
    values : np.array, shape [n_ixs,]
        The array with the ground-truth values.
    ixs : tuple, shape [2, n_ixs]
        The indices at which we want to evaluate the loss (usu
        ally the nonzero indices of the unshifted data matrix).
    Q : np.array, shape [N, k]
        The matrix Q of a latent factor model.
    P : np.array, shape [k, D]
        The matrix P of a latent factor model.
    reg_lambda : float
        The regularization strength

    Returns
    -----
    loss : float
        The loss of the latent factor model.

    """
    mean_sse_loss = np.sum((values - Q.dot(P)[ixs])**2)
    regularization_loss = reg_lambda * (np.sum(np.linalg.norm
    (P, axis=0)**2) + np.sum(np.linalg.norm(Q, axis=1) ** 2))

    return mean_sse_loss + regularization_loss
```

Alternating optimization

In the first step, we will approach the problem via alternating optimization, as learned in the lecture. That is, during each iteration you first update Q while having P fixed and then vice versa.

Task 2: Implement a function that initializes the latent factors Q and P

In [12]:

```
def initialize_Q_P(matrix, k, init='random'):
    """
    Initialize the matrices Q and P for a latent factor model.

    Parameters
    -----
    matrix : sp.spmatrix, shape [N, D]
        The matrix to be factorized.
    k       : int
        The number of latent dimensions.
    init    : str in ['svd', 'random'], default: 'random'
        The initialization strategy. 'svd' means that we
        use SVD to initialize P and Q, 'random' means we initialize
        the entries in P and Q randomly in the interval [
        0, 1).

    Returns
    -----
    Q : np.array, shape [N, k]
        The initialized matrix Q of a latent factor model.

    P : np.array, shape [k, D]
        The initialized matrix P of a latent factor model.
    """
    np.random.seed(0)

    # TODO: Compute Q and P
    N = matrix.shape[0]
    D = matrix.shape[1]
    if(init == 'random'):
        Q = np.random.rand(N,k)
        P = np.random.rand(k,D)
    else:
        u, s, vt = svds(matrix)
        P = u @ s
        Q = vt

    assert Q.shape == (matrix.shape[0], k)
    assert P.shape == (k, matrix.shape[1])
    return Q, P
```

Task 3: Implement the alternating optimization approach

In [13]:

```
def latent_factor_alternating_optimization(M, non_zero_idx, k,
                                           reg_lambda, max_steps=100, init='random',
                                           log_every=1, patience=5, eval_every=1):
    """
        Perform matrix factorization using alternating optimization. Training is done via patience,
        i.e. we stop training after we observe no improvement on the validation loss for a certain
        amount of training steps. We then return the best values for Q and P observed during training.

        Parameters
        -----
        M                : sp.spmatrix, shape [N, D]
                           The input matrix to be factorized.

        non_zero_idx      : np.array, shape [nnz, 2]
                           The indices of the non-zero entries of the un-shifted matrix to be factorized.
                           nnz refers to the number of non-zero entries. Note that this may be different
                           from the number of non-zero entries in the input matrix M, e.g. in the case
                           that all ratings by a user have the same value.

        k                : int
                           The latent factor dimension.

        val_idx          : tuple, shape [2, n_validation]
                           Tuple of the validation set indices. n_validation refers to the size of the
                           validation set.

        val_values       : np.array, shape [n_validation, ]
                           The values in the validation set.

        reg_lambda       : float
                           The regularization strength.

        max_steps        : int, optional, default: 100
                           Maximum number of training steps. Note that we will stop early if we observe
```

no improvement on the validation error
for a specified number of steps
(see "patience" for details).

`init` : str in ['random', 'svd'], default 'random'
The initialization strategy for P and Q. See function `initialize_Q_P` for details.

`log_every` : int, optional, default: 1
Log the training status every X iterations.

`patience` : int, optional, default: 5
Stop training after we observe no improvement of the validation loss for X evaluation iterations (see `eval_every` for details). After we stop training, we restore the best observed values for Q and P (based on the validation loss) and return them.

`eval_every` : int, optional, default: 1
Evaluate the training and validation loss every X steps. If we observe no improvement of the validation error, we decrease our patience by 1, else we reset it to *patience*.

Returns

`best_Q` : np.array, shape [N, k]
Best value for Q (based on validation loss) observed during training

`best_P` : np.array, shape [k, D]
Best value for P (based on validation loss) observed during training

`validation_losses` : list of floats
Validation loss for every evaluation iteration, can be used for plotting the validation loss over time.

`train_losses` : list of floats
Training loss for every evaluation iteration, can be used for plotting the training loss over time.

`converged_after` : int
it - patience*eval every, where it is

*the iteration in which patience hits 0,
or -1 if we hit max_steps before converging.*

```
"""  
  
# TODO: Compute best_Q, best_P, validation_losses, train_losses and converged_after  
  
Q, P = initialize_Q_P(M, k, init)  
best_Q = 0  
best_P = 0  
validation_losses = []  
train_losses = []  
converged_after = 0  
fit_intercept = False  
original_patience = patience  
prev_validation_error = 0  
  
rr_P = Ridge(alpha=reg_lambda)  
rr_Q = Ridge(alpha=reg_lambda)  
  
train_idx = M.nonzero()  
_,_, train_values = sp.find(M)  
  
stop = False  
count_patience = 0  
noImprovement = 0  
  
while(converged_after < max_steps and not stop):  
    rr_P = rr_P.fit(Q, M)  
    P = rr_P.coef_.T  
  
    rr_Q = rr_Q.fit(P.T, M.T)  
    Q = rr_Q.coef_  
  
    if(converged_after % eval_every == 0):  
        validation_error = loss(val_values, val_idx, Q, P,  
reg_lambda)  
        train_error = loss(train_values, train_idx, Q, P,  
reg_lambda)  
        validation_losses.append(validation_error)  
        train_losses.append(train_error)  
  
        if (len(validation_losses) > 1):  
            if(validation_error >= validation_losses[-2]):  
                patience -= 1
```

```

        noImprovement += 1

    else:
        patience = original_patience
        best_P = P
        best_Q = Q
        noImprovement = 0

    if(noImprovement == patience):
        stop = True

    converged_after+=1

    if(converged_after % log_every == 0 and len(validation
_losses) > 1):
        print(f"Iteration {converged_after}, Train Error:
{train_losses[-1]}, Validation Error: {validation_losses[-1]}"
        )

    return best_Q, best_P, validation_losses, train_losses, co
nverged_after

```

Train the latent factor (nothing to do here)

In [14]:

```

Q, P, val_loss, train_loss, converged = latent_factor_alternat
ing_optimization(M_shifted, nonzero_indices,

k=100, val_idx=val_idx,

val_values=val_values_shifted,

reg_lambda=1e-4, init='random',

max_steps=100, patience=10)

```


Iteration 2, Train Error: 112593.71599916997, Validation Error: 244.08298773093534
Iteration 3, Train Error: 109138.30631146854, Validation Error: 241.9162392621014
Iteration 4, Train Error: 108109.30699564128, Validation Error: 240.6777749118207
Iteration 5, Train Error: 107701.7854206022, Validation Error: 240.1017507813578
Iteration 6, Train Error: 107508.04165209441, Validation Error: 239.78316060737086
Iteration 7, Train Error: 107402.53211614337, Validation Error: 239.5725894762025
Iteration 8, Train Error: 107339.0629723255, Validation Error: 239.43108141595647
Iteration 9, Train Error: 107298.0624986813, Validation Error: 239.34599955778714
Iteration 10, Train Error: 107270.08091241548, Validation Error: 239.3094900829268
Iteration 11, Train Error: 107250.03527464424, Validation Error: 239.31364799949782
Iteration 12, Train Error: 107234.98064089604, Validation Error: 239.34954756447107
Iteration 13, Train Error: 107223.14173628515, Validation Error: 239.40775733822056
Iteration 14, Train Error: 107213.4362298555, Validation Error: 239.47939412971326
Iteration 15, Train Error: 107205.20878231204, Validation Error: 239.55702669032163

Plot the validation and training losses over for each iteration (nothing to do here)

In [15]:


```
fig, ax = plt.subplots(1, 2, figsize=[10, 5])
fig.suptitle("Alternating optimization, k=100")

ax[0].plot(train_loss[1::])
ax[0].set_title('Training loss')
plt.xlabel("Training iteration")
plt.ylabel("Loss")

ax[1].plot(val_loss[1::])
ax[1].set_title('Validation loss')
plt.xlabel("Training iteration")
plt.ylabel("Loss")

plt.show()
```

Alternating optimization, k=100



In []:

Exporting the results to PDF

Once you complete the assignments, export the entire notebook as PDF and attach it to your homework solutions. The best way of doing that is

1. Run all the cells of the notebook.
2. Export/download the notebook as PDF (File -> Download as -> PDF via LaTeX (.pdf)).
3. Concatenate your solutions for other tasks with the output of Step 2. On linux, you can use `pdffunite`, there are similar tools for other platforms, too. You can only upload a single PDF file to Moodle.

Make sure you are using `nbconvert` version 5.5 or later by running `jupyter nbconvert --version`. Older versions clip lines that exceed page width, which makes your code harder to grade.

In [1]:

```
import matplotlib.pyplot as plt
import numpy as np
from PIL import Image
from sklearn.datasets import load_sample_image

%matplotlib inline

def compare_images(img, img_compressed, k):
    """Show the compressed and uncompressed image side by side
    .
    """

    fig, axes = plt.subplots(1, 2, figsize=(16, 12))
    axes[0].set_axis_off()
    if isinstance(k, str):
        axes[0].set_title(k)
    else:
        axes[0].set_title(f"Compressed to {k} colors")
    axes[0].imshow(img_compressed)
    axes[1].set_axis_off()
    axes[1].set_title("Original")
    axes[1].imshow(img)
```

K-Means

In this first section you will implement the image compression algorithm from Bishop, chapter 9.1.1. Take an RGB image $X \in \mathbb{R}^{h \times w \times 3}$ and interpret it as a data matrix $X \in \mathbb{R}^{N \times 3}$. Now apply k -means clustering to find k colors that describe the image well and replace each pixel with its associated cluster.

In [2]:

```
# Alternatively try china.jpg
X = load_sample_image("flower.jpg")

# or load your own image
# X = np.array(Image.open("/home/user/path/to/some.jpg"))
```

In [3]:

```
def computeDistance(point, mu):
    d = [np.linalg.norm(point-centroid) for centroid in mu]
    return d

def kmeans(X, k):
    """Compute a k-means clustering for the data X.

    Parameters
    -----
    X : np.array of size N x D
        where N is the number of samples and D is the data dimensionality
    k : int
        Number of clusters

    Returns
    -----
    mu : np.array of size k x D
        Cluster centers
    z : np.array of size N
        Cluster indicators, i.e. a number in 0..k - 1, for each data point in X
    """

    # TODO: Compute mu and z

    # Initialize centroids
    mu = X[k]
```

```

mu = x[:k]
z = []
distances = []
clusters = [[] for _ in range(k)]
convergence = False
prevZ = []

while (not convergence):
    # Compute distance of each point to the centroids and
    # assign each point to a cluster
    prevZ = z
    for point in X:
        d = computeDistance(point, mu)
        cluster = np.argmin(d)
        z.append(cluster)
        clusters[cluster].append(point)

    # Compute the new centroids
    for i in range(k):
        mu[i] = np.mean(clusters[i])

    if(prevZ == z):
        convergence = True

return mu, z

```

In [4]:

```


# Cluster the color values
k = 5
mu, z = kmeans(X.reshape((-1, 3)), k)

# Replace each pixel with its cluster color
X_compressed = mu[z].reshape(X.shape).astype(np.uint8)


# Show the images side by side
compare_images(X, X_compressed, k)

```

Compressed to 5 colors



Original



Gaussian Mixture Models & EM

Now you will repeat the same exercise with GMMs.

In [5]:

```
def gmm_log_probability(X, pi, mu, sigma):
    """Compute the joint log-probabilities for each data point
    and component.

    Parameters
    -----
    X : np.array of size N x D
        where N is the number of samples and D is the data dim
    ensionality
    pi : np.array of size k
        Prior weight of each component
    mu : np.array of size k x D
        Mean vectors of the k Gaussian component distributions
    sigma : np.array of size k x D x D
        Covariance matrices of the k Gaussian component distri
    butions

    Returns
    -----
    P : np.array of shape N x k
        P[i, j] is the joint log-probability of data point i u
    nder component j
    """

    # TODO: Compute P
    K = len(pi)
    P = np.log(np.sum([pi[i]*np.random.multivariate_normal(mu[
i],cov[i]).pdf(X[0]) for i in range(K)]))
    print(P)
    log_likelihoods = []
    log_likelihoods.append(np.log(np.sum([k*multivariate_norma
l(mu[i],sigma[j]).pdf(X) for k,i,j in zip(pi,range(len(mu)),ra
nge(len(sigma)))])))
    print(log_likelihoods)
    return P

def em(X, k, tol=0.001):
    """Fit a Gaussian mixture model with k components to X.
```

Parameters

X : np.array of size N x D
where N is the number of samples and D is the data dimensionality
k : int
Number of clusters
tol : float
Converge when the increase in the mean of the expected joint log-likelihood
is lower than this

The algorithm should stop when the relative improvement in the optimization objective is less than rtol.

Returns

pi : np.array of size k
Prior weight of each component
mu : np.array of size k x D
Mean vectors of the k Gaussian component distributions
sigma : np.array of size k x D x D
Covariance matrices of the k Gaussian component distributions
"""

return pi, mu, sigma

In [6]:

```
# Fit the GMM
k = 5
pi, mu, sigma = em(X.reshape((-1, 3)), k)

# Determine the most likely cluster of each pixel
log_p = gmm_log_probability(X.reshape((-1, 3)), pi, mu, sigma)
z = log_p.argmax(axis=1)

# Replace each pixel with its cluster mean
X_compressed = mu[z].reshape(X.shape).astype(np.uint8)

# Show the images side by side
compare_images(X, X_compressed, k)
```

```
-----
-----
NameError
Traceback (most recent call last)
<ipython-input-6-5c2c54e53fa0> in <module>
      1 # Fit the GMM
      2 k = 5
----> 3 pi, mu, sigma = em(X.reshape((-1, 3)), k)
      4
      5 # Determine the most likely cluster of each pixel

<ipython-input-5-1761a0cee8fe> in em(X, k, tol)
     56
     57
---> 58     return pi, mu, sigma

NameError: name 'pi' is not defined
```

Sampling Unseen Datapoints

You have trained a generative model which allows you to sample from the learned distribution. In this section, you sample new images.

In []:

```
def gmm_sample(N, pi, mu, sigma):
    """Sample N data points from a Gaussian mixture model.

    Parameters
    -----
    N : int
        Number of data points to sample
    pi : np.array of size k
        Prior weight of each component
    mu : np.array of size k x D
        Mean vectors of the k Gaussian component distributions
    sigma : np.array of size k x D x D
        Covariance matrices of the k Gaussian component distributions

    Returns
    -----
    X : np.array of shape N x D
    """

    # TODO: Sample X

    return X
```

In []:

```
# Sample pixels and reshape them into the size of the original image
X_sampled = gmm_sample(np.prod(X.shape[:-1]), pi, mu, sigma).r
eshape(X.shape).astype(np.uint8)

# Compare the original and the sampled image
compare_images(X, X_sampled, "Sampled")
```

Explain what you see in the generated images. (1-3 sentences)